

1. Launch the OpenShift cluster

You are using **Code-Ready Containers**, a single-server OpenShift cluster that Red Hat makes available for free. It's the successor of minishift.

To start the cluster, run **crc start**. After about five minutes, the cluster is up and running.

Copy the last few output lines into a file:

Started the OpenShift cluster.

The server is accessible via web console at:
<https://console-openshift-console.apps-crc.testing>

Log in as administrator:
Username: kubeadmin
Password: vNd8o-Ryg8J-hYLSP-AoGpU

Log in as user:
Username: developer
Password: developer

Use the 'oc' command line interface:
\$ eval \$(crc oc-env)
\$ oc login -u developer https://api.crc.testing:6443

The kubeadmin password is particularly important.

2. Deploy an app from an image

First contact with OpenShift

```
$ oc login -u developer
error: couldn't get https://192.168.49.2:8443/.well-known/oauth-authorization-server: unexpected
response status 403
```

Before you can use the kubectl or oc command, you need to set the kubeconfig context to CRC.

```
$ oc config get-contexts
```

CURRENT	NAME	CLUSTER	AUTHINFO	NAMESPACE
	/api-crc-testing:6443/developer	api-crc-testing:6443	developer/api-crc-testing:6443	
	crc-admin	api-crc-testing:6443	kubeadmin	default
	crc-developer	api-crc-testing:6443	developer	default
	default/api-...ng:6443/kubeadmin	api-crc-testing:6443	kubeadmin/api-crc-testing:6443	default
	devproject/api-...ng:6443/developer	api-crc-testing:6443	developer/api-crc-testing:6443	devproject
	devproject/api-...ng:6443/kubeadmin	api-crc-testing:6443	kubeadmin/api-crc-testing:6443	devproject
*	minikube	minikube	minikube	default

Currently, the kubeconfig file (`$HOME/.kube/config`) uses minikube as context, since this is what you have been using so far. Change it to one of the crc contexts. It doesn't matter which.

```
$ oc config use-context crc-admin
Switched to context "crc-admin".
$ oc login -u developer
Password: developer
Login successful.
```

You don't have any projects. You can try to create a new project, by running

```
oc new-project <projectname>
```

```
$ oc whoami
```

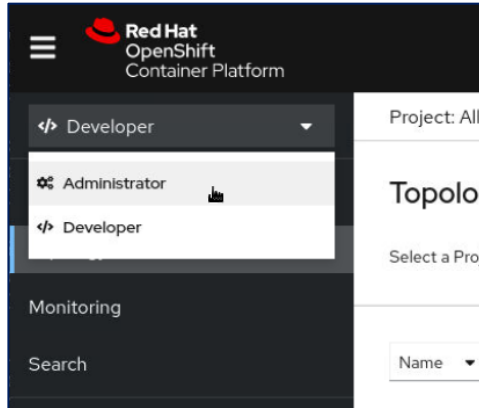
developer

Create a project:

```
$ oc new-project devproject
```

Projects are OpenShift's method to group applications. They have access control and quotas for developers or teams. Projects are based on Kubernetes namespaces.

Access the OpenShift web console at <https://console-openshift-console.apps-crc.testing>. Log on as developer, password developer.



The Topology page is shown, with the project you just created.

If you forgot to create the project, do it from the console: Open the **Project** menu, then click on **create a Project**. Only the project's name has to be provided.

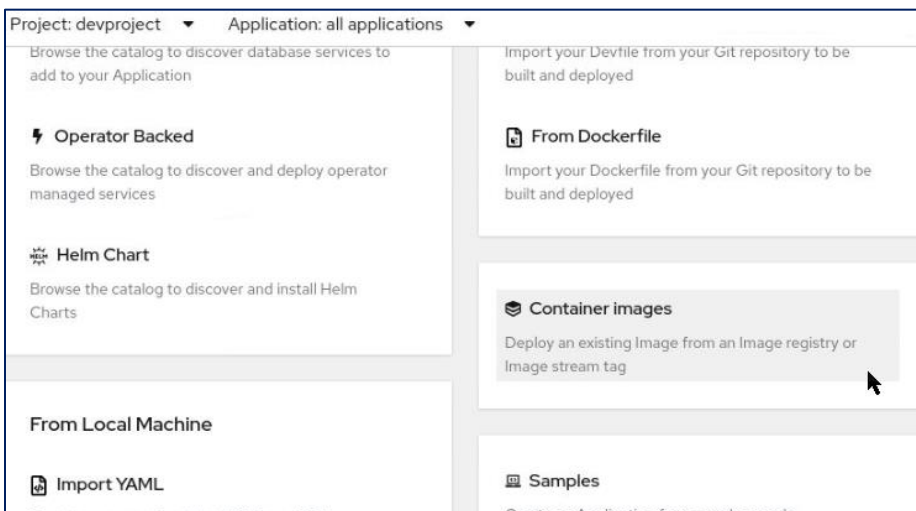
Your perspective is *Developer*. Switch to the *Administrator* perspective (screenshot on the left) and note how the menu changes. Go back to the *Developer* perspective.

The Administrator Perspective consists of: Operators, Workloads, Networking, Storage, Builds, User Mgmt and Administration menus.

The Developer Perspective includes: Topology, Monitoring, Search, Builds, Helm, Project, ConfigMaps, Secrets.

Deploy an app

In the developer perspective, enter **+Add**. This section allows you to create applications from various sources and using various methods. View the available options. Find the **Container images** tile and click it. You will now deploy an app from an image.



Under *Image name from external registry*, enter image name `docker.io/openshiftroadshow/parksmat-katacoda:1.2.0`. The image will be validated.

If you like, you can select a different icon to identify your application.

Scroll down. Two names were set automatically; the app is named `parksmat-katacod-app`, and the name `parksmat-katacoda` will be used as a basis for naming various application components.

Under **Resources**, you have the choice of **Deployment** or **DeploymentConfig**. Select **Deployment** for now.

Under **Advanced Options**, the **route** box is checked. Since you will create the route later, **uncheck it**.

There are advanced options for **Health Checks**, **Deployment**, **Scaling**, **Resource Limits** and **Labels**. Open them all. Most of them were covered to some extent during the Kubernetes course module; try to remember their meaning. Feel free to set some options.

Click **Create**. The browser goes back to *Topology*, where this picture is shown:



The outer white ring is the app, the inner blue ring is a Deployment.

Switch to a terminal window and list your resources

```
$ oc get all
NAME                                READY   STATUS    RESTARTS   AGE
pod/parksmmap-katacoda-6849877f77-xs6s9  1/1     Running   0           118s

NAME                                TYPE          CLUSTER-IP    EXTERNAL-IP  PORT(S)    AGE
service/parksmmap-katacoda            ClusterIP     10.217.5.149  <none>       8080/TCP   118s

NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/parksmmap-katacoda  1/1     1             1           118s

NAME                                DESIRED   CURRENT   READY   AGE
replicaset.apps/parksmmap-katacoda-6849877f77  1         1         1       118s

NAME                                IMAGE REPOSITORY
TAGS      UPDATED
imagestream.image.openshift.io/parksmmap-katacoda  default-route-openshift-image-registry.apps-
crc.testing/devproject/parksmmap-katacoda  1.2.0  About a minute ago
```

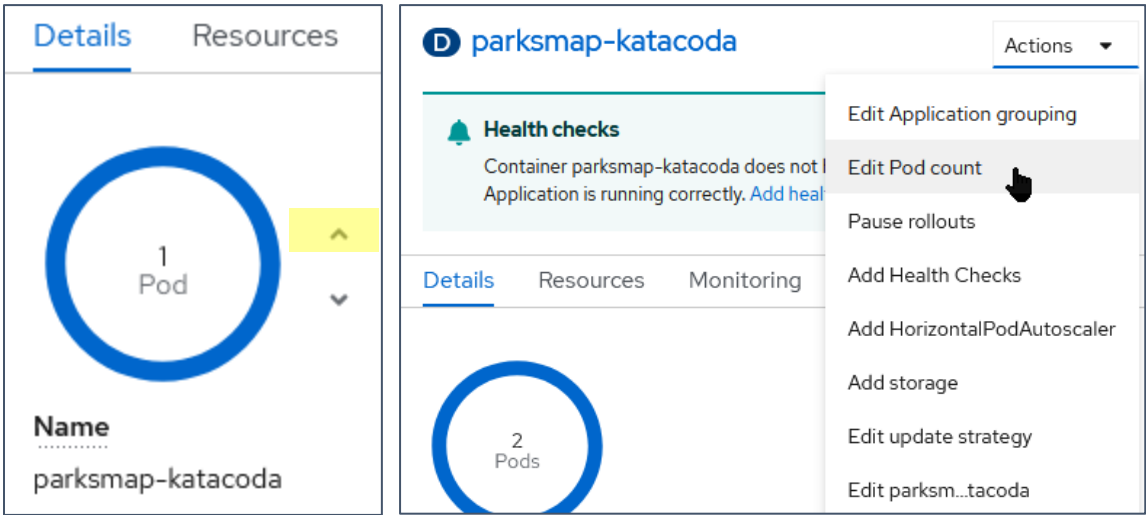
There is a deployment, a replicaset, a pod, a service and the OpenShift-specific resource imagestream.

View details in the browser

Click on the blue ring to open the Deployment menu. The Resources tab lists the pod and the service.


Open the *Details* tab. This seems to be a graphical representation of `oc describe deployment`. Click on the dash-underlined headers such as **Update strategy** to see a short explanation of their meaning.

Scale the application up by clicking on the up arrow, or Actions → Edit Pod Count



Go back to the **resources** tab. There are two pods now. Click on a pod.

Lots of pod details are shown, including metrics. The **YAML** tab displays the pod's manifest and allows you to download it.



Select action **delete pod**. It will be deleted, but since the replica count is 2, a new pod is created immediately.

Scale back down to 1. This removes one of the two pods.

Create a route

Switch to the administrator perspective. Select **myproject**. In the left menu, select **Networking→Routes**. Create a route and set a name, select the parksmat service and the 8080→8080 port mapping; leave **secure route** unchecked.

After clicking Create, you are in the route's details page. The route's URL should be similar to <http://parkmap-route-devproject.apps-crc.testing>.

Go back to **developer** perspective and select Topology. The route is indicated by this icon:



Click on the route icon to open the app. Currently, it only shows a world map that can be moved and zoomed.

Building from source code

You will add a second deployment to your app which provides locations of national parks around the world. This time, you will use S2i to generate the app.

Ensure that you are in the developer perspective.

Click **+Add** in the menu. Click **Catalog→All Services**.

Notice the type of objects in the catalog:

- Builder Images**; S2i-enabled container images for building source code for with a particular language or framework.

- Helm Charts**

- Templates**; sample applications that you can tweak to build your own.

- Devfiles**; YAML code that describes a development environment.

- Not shown here: Operator backed**; infrastructure services managed by Kubernetes controllers.

Select **Languages** on the left menu, then **Python**. There are three options; click on the builder image **Python**, read the description, then **Create App**.

Enter **Git Repo URL** <https://github.com/openshift-roadshow/nationalparks-katacoda>. OpenShift performs a quick validation: Is this a valid repo, is it accessible.

Leave all other options unchanged and click **Create**. This sends you back to the **Topology** page.

The parksmat app now has two deployments, one being the Python code.



The light blue ring and the hourglass icon indicate that the application is waiting to build. The Github icon links to the source code's repo. The hourglass will eventually change to an icon indicating that a build is in progress.



Click on the Python ring and select **Resources** → **Builds**. The build might be running but will eventually complete. **View logs** shows the build messages. The last message should be *Push successful*.

Back to the **Topology** page: The Python deployment has a check mark now, meaning that the build is complete. The ring is a darker blue, also indicating that the application is ready. Click on the route icon of the left deployment to see locations of national parks.

Explore resources on the command line

Go to a terminal and list all resources.

```
$ oc get all
NAME                                     READY   STATUS    RESTARTS   AGE
pod/nationalparks-katacoda-1-build      0/1     Completed 0           7m54s
...

NAME                                     DESIRED   CURRENT   READY   AGE
replicaset.apps/nationalparks-katacoda-74ff6bdbbf 1         1         1       2m44s
replicaset.apps/nationalparks-katacoda-85898c9f8f 0         0         0       7m54s
replicaset.apps/parksmap-katacoda-6849877f77      2         2         2       40m

NAME                                     TYPE     FROM     LATEST
buildconfig.build.openshift.io/nationalparks-katacoda Source   Git      1

NAME                                     TYPE     FROM     STATUS    STARTED
DURATION
build.build.openshift.io/nationalparks-katacoda-1 Source   Git@43842c2 Complete  7 minutes ago
5m10s
...
NAME                                     HOST/PORT
PATH   SERVICES          PORT    TERMINATION  WILDCARD
route.route.openshift.io/nationalparks-katacoda nationalparks-katacoda-devproject.apps-crc.testing
nationalparks-katacoda 8080-tcp      None
```

```
route.route.openshift.io/parkmap-route
parkmap-katacoda      8080-tcp
```

```
parkmap-route-devproject.apps-crc.testing
None
```

Many more resources exist now.

One of the pods was used for building the app and is now complete. It is associated with the Build resource:

```
$ oc describe pod/nationalparks-katacoda-1-build | grep Controlled
Controlled By: Build/nationalparks-katacoda-1
```

List the details of the build. A few points worth noting: The name and some details of the build config are included, the reason why the build was started, and there is a long list of events.

```
$ oc describe build nationalparks-katacoda-1 | more
```

Also display the details of the build config. They should match what you say in the build.

There are two versions of the same replicaset because the corresponding deployment was upgraded:

```
$ oc rollout history deploy nationalparks-katacoda
deployment.apps/nationalparks-katacoda
REVISION  CHANGE-CAUSE
1          <none>
2          <none>
```

To understand why there are two revisions, display their details. The image was updated from revision 1 to 2.

```
$ oc rollout history deploy nationalparks-katacoda --revision=1
$ oc rollout history deploy nationalparks-katacoda --revision=2
```

3. Deploy a second app from source

Create a second project from the command line or the webconsole.

```
$ oc new-project myproject
Now using project "myproject" on server "https://openshift:6443".
```

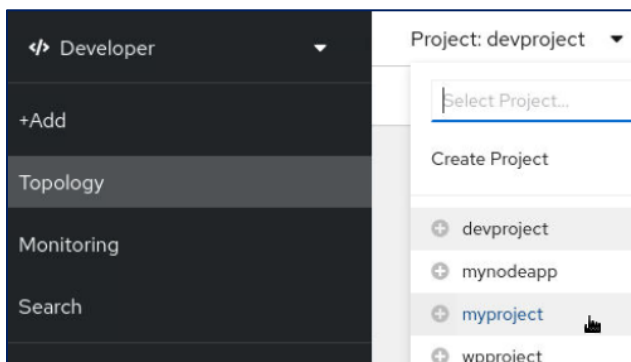
You can add applications to this project with the 'new-app' command. For example, try:

```
oc new-app rails-postgresql-example
```

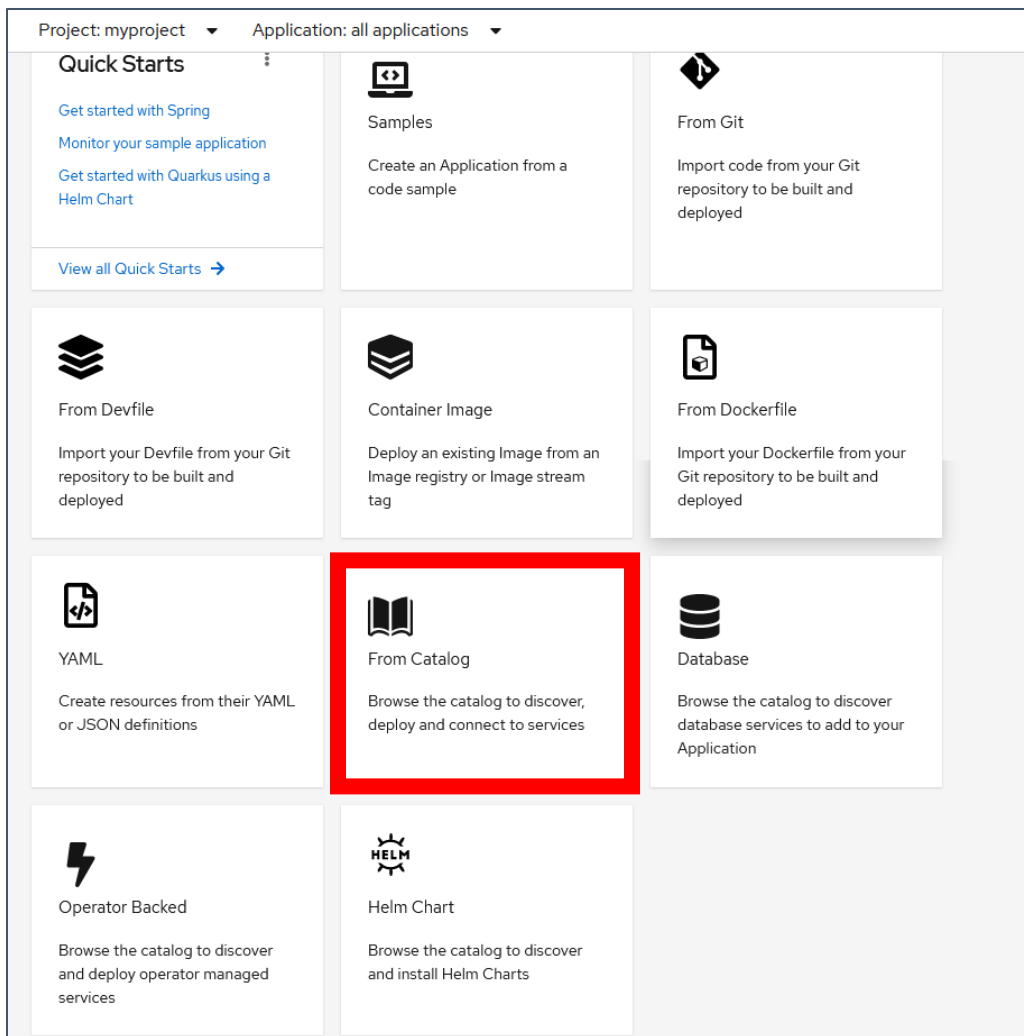
to build a new example application in Ruby. Or use `kubectl` to deploy a simple Kubernetes application:

```
kubectl create deployment hello-node --image=k8s.gcr.io/serve_hostname
```

In the console, ensure you are developer, go to the **Topology** page and select the new project from the drop-down menu.



Click on **myproject**. The Topology page is empty. Go to **+Add** and select **From Catalog**.



Select **Language** → **Python**. Click on the **Python** tile.

The Python page describes the Python builder image. **Create Application**.

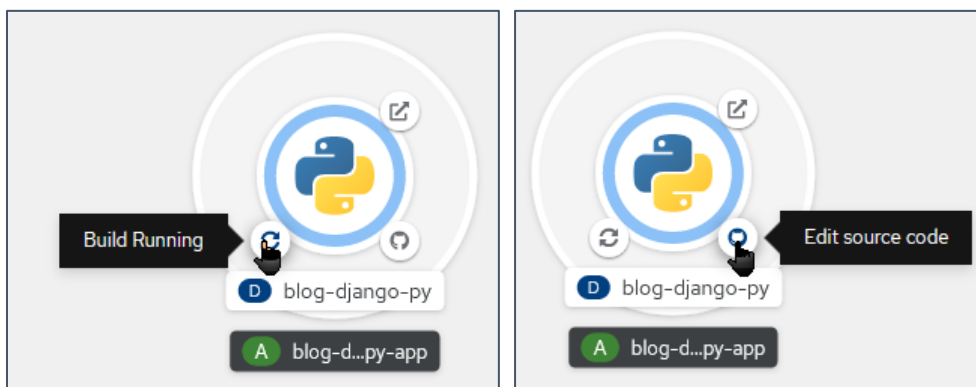
Git Repo URL: Enter <https://github.com/openshift-katacoda/blog-django-py>.

General: see app name and name.

Leave **Route** checkbox checked.

Other **Advanced Settings**: Routing, Health Checks, Build config, Deployment, Scaling, Resource limits.

Leave everything unchanged and **Create**.



The ring color goes from white (build pending) to light blue (build running) to a darker blue (app running)

Clicking on the **Build Running** icon shows the logs. **Edit Source Code** links to the Git repo.

Delete app

Switch to the command line.

Show all resources. Assuming this is your only currently active app, you can simply list all resources:

```
$ oc get all -o name
pod/blog-django-py-1-build
pod/blog-django-py-6d77767b8c-ctj8m
service/blog-django-py
deployment.apps/blog-django-py
replicaset.apps/blog-django-py-6d77767b8c
replicaset.apps/blog-django-py-7fc8bdfb
buildconfig.build.openshift.io/blog-django-py
build.build.openshift.io/blog-django-py-1
imagestream.image.openshift.io/blog-django-py
route.route.openshift.io/blog-django-py
```

Notice that the app comprises pods, a service, a deployment, replicaset, a buildconfig, a build, an imagestream and the route.

In case several apps are deployed, use a label to filter the right app.

```
$ oc get all -o name -l app=
```

Display details of the route:

```
$ oc describe route blog-django-py
Name:                blog-django-py
Namespace:           myproject
Created:             8 minutes ago
Labels:              app=blog-django-py
                    app.kubernetes.io/component=blog-django-py
                    app.kubernetes.io/instance=blog-django-py
                    app.kubernetes.io/name=python
                    app.kubernetes.io/part-of=blog-django-py-app
                    app.openshift.io/runtime=python
                    app.openshift.io/runtime-version=3.8-ubi7
Annotations:         openshift.io/host.generated=true
Requested Host:      blog-django-py-myproject.2886795282-80-hazel04.environments.katacoda.com
                    exposed on router default (host apps-crc.testing) 8 minutes ago
Path:                <none>
TLS Termination:     <none>
Insecure Policy:     <none>
Endpoint Port:       8080-tcp

Service:            blog-django-py
Weight:             100 (100%)
Endpoints:          10.217.1.13:8080
```

The label `app=blog-django-py` was set by the web console. List all resources with this label.

```
$ oc get all --selector app=blog-django-py -o name
$ oc delete all -o name --selector app=blog-django-py
pod/blog-django-py-7fc8bdfb-6jjmw
service/blog-django-py
deployment.apps/blog-django-py
replicaset.apps/blog-django-py-7fc8bdfb
buildconfig.build.openshift.io/blog-django-py
build.build.openshift.io/blog-django-py-1
imagestream.image.openshift.io/blog-django-py
route.route.openshift.io/blog-django-py
```

The label in this example is not universally correct. Always use `oc describe` to verify what labels have been applied. Before deleting, check with `oc get all --selector`.

Deploying using the command line

Ensure you are in `devproject`.

```
$ oc projects
$ oc project
$ oc project devproject
```

The `new-app` command below uses `IMAGE~CODE` syntax. It uses builder image `python:latest` and clones the source code from Github.

```
$ oc new-app python:latest~https://github.com/openshift-katacoda/blog-django-py
--> Found image f742fc9 (5 months old) in image stream "openshift/python" under tag "latest" for
"python:latest"
...
--> Creating resources ...
...
--> Success
    Build scheduled, use 'oc logs -f buildconfig/blog-django-py' to track its progress.
    Application is not exposed. You can expose services to the outside world by executing one or more of
the commands below:
    'oc expose service/blog-django-py'
    Run 'oc status' to view your app.
```

The application is not exposed, which means there is no route. In contrast to the web console build process, the `new-app` command does not create a route by default..

```
$ oc logs bc/blog-django-py --follow
```

`bc` is the abbreviation for `buildconfig`. This shows the `s2i` output including buildah output. It will take several minutes to complete. You can interrupt and continue with the next commands. The following command shows the status of services, deployment configs, build configurations, and deployments in the current project. It looks like the Django blog is still building. The national parks app consists of two deployments, which are in progress.

```
$ oc status
```

In project `devproject` on server `https://api.crc.testing:6443`

```
svc/blog-django-py - 10.217.5.48:8080
  deployment/blog-django-py deploys istag/blog-django-py:latest <-
    bc/blog-django-py source builds https://github.com/openshift-katacoda/blog-django-py on
openshift/python:latest
    build #1 running for 2 minutes - 9c4f338: Merge pull request #9 from mixpix3ls/master (ryan
jarvinen <ryan.jarvinen@gmail.com>)
    deployment #1 running for 2 minutes - 0/1 pods growing to 1
...
3 infos identified, use 'oc status --suggest' to see details.
$ oc get all
```

This is too much information. Use a label to filter for the Django app.

```
$ oc get all -l app=blog-django-py
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/blog-django-py	ClusterIP	10.217.5.48	<none>	8080/TCP	10m

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/blog-django-py	1/1	1	1	10m

NAME	TYPE	FROM	LATEST
buildconfig.build.openshift.io/blog-django-py	Source	Git	1

NAME	TYPE	FROM	STATUS	STARTED	DURATION
build.build.openshift.io/blog-django-py-1	Source	Git@9c4f338	Complete	10 minutes ago	4m39s

NAME	IMAGE	REPOSITORY
TAGS	UPDATED	

```
imagestream.image.openshift.io/blog-django-py default-route-openshift-image-registry.apps-
crc.testing/devproject/blog-django-py latest 6 minutes ago
```

Replicasets and pods have a different label.

```
$ oc get all -l deployment=blog-django-py
```

NAME	READY	STATUS	RESTARTS	AGE
pod/blog-django-py-b87967fbd-trmp9	1/1	Running	0	6m32s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/blog-django-py-78b7478997	0	0	0	11m
replicaset.apps/blog-django-py-b87967fbd	1	1	1	6m32s

Use the `expose` command to create a route.

```
$ oc expose service blog-django-py
```

```
route.route.openshift.io/blog-django-py exposed
```

```
$ oc get route
```

NAME	HOST/PORT	PATH
SERVICES	PORT	TERMINATION WILDCARD
blog-django-py	blog-django-py-myproject.2886795286-80-elsy07.environments.katacoda.com	blog-
django-py	8080-tcp	None
...		

Check the Topology page. The Django deployment is there but has no app. This is so because it lacks a label. Add the required label, then check the topology again.

```
$ oc label deploy blog-django-py app.kubernetes.io/part-of=blog-django-app
```

Try the route icon.

Trigger a new build

Go back to the command line and manually trigger a build. This command requires a BuildConfig as argument.

```
$ oc get bc
```

NAME	TYPE	FROM	LATEST
blog-django-py	Source	Git	1
nationalparks-katacoda	Source	Git	1

```
$ oc start-build blog-django-py
```

```
build.build.openshift.io/blog-django-py-2 started
```

Follow the build.

```
$ oc get build --watch
```

NAME	TYPE	FROM	STATUS	STARTED	DURATION
nationalparks-katacoda-1	Source	Git@43842c2	Complete	2 hours ago	5m10s
blog-django-py-1	Source	Git@9c4f338	Complete	28 minutes ago	4m39s
blog-django-py-2	Source	Git@9c4f338	Running	15 seconds ago	

The logs have more detail.

```
$ oc logs buildconfig/blog-django-py
```

```
...
```

```
$ oc describe buildconfig blog-django-py
```

```
...
```

```
URL: https://github.com/openshift-katacoda/blog-django-py
```

```
...
```

```
Webhook GitHub:
```

```
URL:
```

```
https://openshift:6443/apis/build.openshift.io/v1/namespaces/myproject/buildconfigs/blog-
django-py/webhooks/<secret>/github
```

```
Webhook Generic:
```

```
URL:
```

```
https://openshift:6443/apis/build.openshift.io/v1/namespaces/myproject/buildconfigs/blog-
django-py/webhooks/<secret>/generic
```

```
AllowEnv: false
```

Webhooks are used to detect that a new build must be performed, e.g. automatically when updated source is pushed to repo. This is possible for repos that you own.

4. Build locally

Use local source code to build an image. Start by cloning the Django blog code.

```
$ git clone https://github.com/openshift-katacoda/blog-django-py
$ cd blog-django-py
```

Make a change to the code. The default color is red. You add an environment variable to the S2i configuration.

```
$ echo 'BLOG_BANNER_COLOR=blue' >> .s2i/environment
```

New build:

```
$ oc start-build blog-django-py --from-dir=. --wait
```

Note the `from-dir` option. This build uses the local directory instead of the git repo. S2i calls this a *binary build*, although this term might be a bit misleading for a Python program.

Also check progress on the web console under the Builds menu: **Builds** → **blog-django-py** → **Builds tab**, then select the running build and the **logs** tab.

Accessing the app after the build shows that the banner has turned from red to blue.

To revert to the git source, run `oc start-build blog-django-py`

Since you are impatient, cancel the ongoing build: `oc cancel-build blog-django-py-NUMBER`

List the deployment's update history with `oc rollout history deploy blog-django-py`. Each build causes a new rollout.

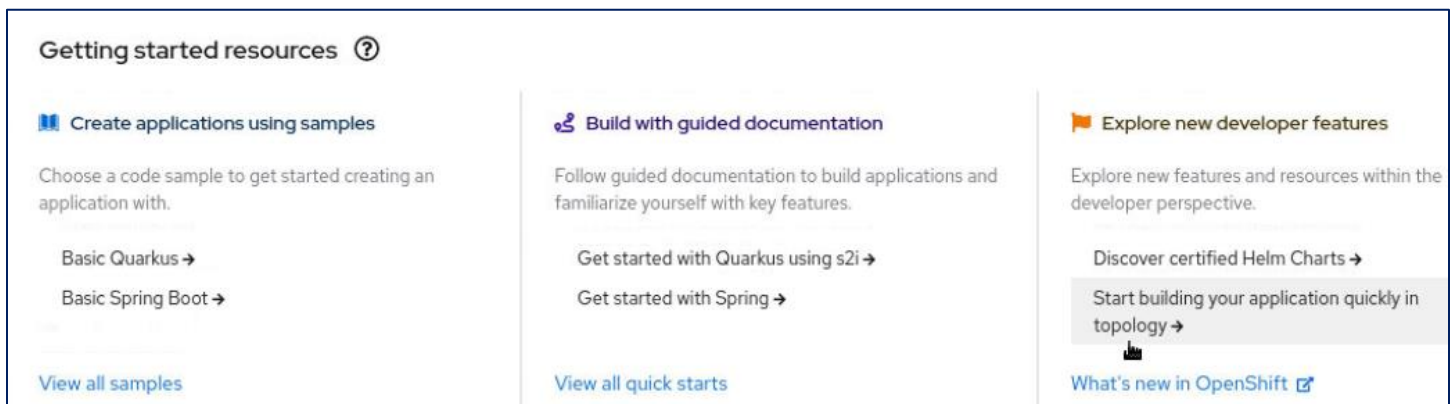
5. Optionally, build the Nodejs app

This is the same nodejs application that you used when building Docker images using Docker and S2i. You will clone the code from Github.

Log on to the web console as developer. Use the developer's perspective.

Go to the **+Add** section. Explore two ways of launching a Github-based build. Find the **Git Repository** tile and click on **From Git**. Explore the options, but don't fill any fields and go back to **+Add**.

Under **Getting started resources**, and **Explore new developer features**, select **Start building your application quickly in topology**.



An input field "Add to Project" is displayed. Enter the text **node**. You will be presented with a list of NodeJS images. Select **Nodejs Builder Images**, then click the blue **Create Application** button.

Select Builder Image Version **14-ubi8**. This is nodejs version 14 on RHEL 8.

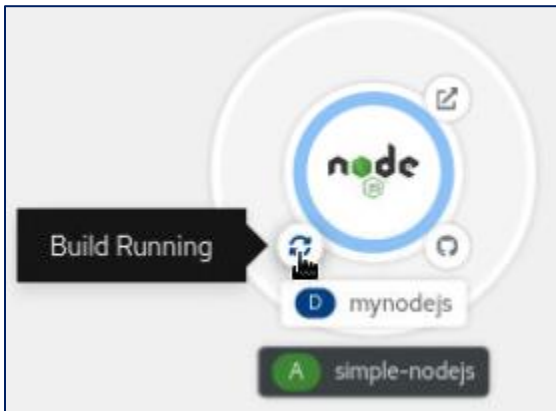
Enter the Git repo <https://github.com/berndbausch/openshift-nodejs>.

Enter **Application Name** *simple-nodejs* and **Name** *mynodejs*. Your BuildConfig will be named *mynodejs*, and other elements will contain *mynodejs* as part of their names.

Scroll down. Leave the **Create a route** box checked, but open the advanced route options and enter **awesome-site.apps-crc.testing** in the hostname field. To be reachable in this CRC environment, all routes must use *apps-crc.testing* as their domain component.

Click the blue **Create** box.

This sends you back to the Topology screen, where you see your app. Click on the build icon to see live log messages from the build.



Eventually, the build will complete. Click on the route icon to see the app in action.

Switch between graphical and list mode by clicking the **Graph View** and **List View** buttons:



Go to the Graph View. Initially, your Deployment is marked with a light blue ring, which will eventually turn into a darker blue. When this is the case, the app is ready. If your mouse has a scroll wheel, you can use it to zoom in and out (there are also zoom buttons).



Light blue means that the app is currently building. This is also symbolized by the small icon in the lower left. This icon turns into a checkmark at the end of the build. Notice the Github icon; clicking it opens the code's Github repo.

Hover your mouse over the dark blue ring. How many pods are running?

Explore the interface by clicking on various elements of this graph, but don't spend too much time trying to understand what they mean.

When you are happy with your exploration, go to **Builds** in the left menu. Your BuildConfigs (labeled with a blue BC) are listed here. Click on the *mynodejs* BuildConfig.

The **Details** tab shows a description of the BuildConfig object, including the source location, the build image, the output image, run policy and triggers. Apart from the source location and the build image, these are default values.

Feel free to click the dash-underlined details for a short description of their meaning.

Go to the **YAML** tab to see the manifest. Download it; it will be saved to **buildconfig-mynodejs.yaml** in your Downloads directory.

Go to the **Builds** tab. It lists the builds that have been performed so far.

Back to the command line, check what project you are in.