

Manifests are under `classfiles/os-nw`. Do a `git pull` first.

Cluster Network Operator

```
$ oc get clusteroperator
```

List of cluster operators incl network

This shows a healthy cluster network operator:

```
$ oc get clusteroperator/network
```

NAME	VERSION	AVAILABLE	PROGRESSING	DEGRADED	SINCE
network	4.8.12	True	False	False	20d

The structure: It's a deployment with one pod.

```
$ oc get -n openshift-network-operator deployment/network-operator
```

```
$ oc get all -n openshift-network-operator
```

Warnings or errors in the log

```
$ oc logs deploy/network-operator -n openshift-network-operator |grep '^[WE]'
```

Cluster network config

Configuration is contained in a customresource:

```
$ oc get networks/cluster -o yaml
```

```
apiVersion: config.openshift.io/v1
```

```
kind: Network
```

```
metadata:
```

```
  creationTimestamp: "2021-09-23T06:59:03Z"
```

```
  generation: 2
```

```
  name: cluster
```

```
  resourceVersion: "2988"
```

```
  uid: ae1c3f58-bd30-4fc1-9d8c-ca07b2be07ae
```

```
spec:      ← As configured
```

```
  clusterNetwork:      ← Pod addresses
```

```
    - cidr: 10.217.0.0/22
```

```
      hostPrefix: 23      ← prefix for each node
```

```
  externalIP:
```

```
    policy: {}
```

```
  networkType: OpenShiftSDN
```

```
  serviceNetwork:      ← Services
```

```
    - 10.217.4.0/23
```

```
status:      ← Actual state
```

```
  clusterNetwork:
```

```
    - cidr: 10.217.0.0/22
```

```
      hostPrefix: 23
```

```
  clusterNetworkMTU: 1400
```

```
  networkType: OpenShiftSDN
```

```
  serviceNetwork:
```

```
    - 10.217.4.0/23
```

Configure the NodePort range

This is an example of configuring the CNO. A NodePort service can use ports between 30000 and 32766 by default. You can increase that range.

First confirm that the highest nodeport is 32767.

Log on to the console as *developer* and switch to the **administrator** perspective. Go to **Networking**→**Services** and create a service. You will be given a YAML file; add the *NodePort* type and a NodePort value greater than 32767:

```
spec:
  selector:
    app: MyApp
  type: NodePort
  ports:
    - protocol: TCP
      port: 80
      nodeport: 35000
      targetPort: 9376
```

When you create this service, you will see that the actual nodeport value is set to less than 32767.

To increase the nodeport range, apply the following manifest `cno-nodeport.yaml` as kubeadmin.

```
apiVersion: config.openshift.io/v1
kind: Network
metadata:
  name: cluster
spec:
  serviceNodePortRange: "30000-35000"
```

It will take a minute or two before this change is effective.

Go back to the web console. Delete the previous service and create a new one. This time, 35000 should be accepted.

In `cno-nodeport.yaml`, reset the range to 30000-32767 and apply the manifest again.

DNS Operator

```
$ oc project openshift-dns-operator
$ oc get all
```

There is a service.

```
$ oc describe dns.operator/default
```

status.conditions tells why the DNS operator is degraded, progressing or available

To see the DNS operator's daemon set:

```
$ oc get ds -n openshift-dns
```

Ingress Operator - access logs

The operator is configured by the `ingresscontroller/default` resource in the `openshift-ingress-operator` namespace.

Configuration options: Certificates, replicas, access logs, endpointPublishingStrategy

Ingress is implemented by a deployment in namespace `openshift-ingress`.

Explore the resources:

```
$ oc describe ingresscontroller/default -n openshift-ingress-operator
$ oc get all -n openshift-ingress
```

Feel free to [describe](#) the resources you see.

Configure access logging

As an example how you can configure the OpenShift ingress operator, enable access logging. This will enable you to see accesses to Ingresses and Routes.

Use this manifest named `ingress-log.yaml`.

```
apiVersion: operator.openshift.io/v1
kind: IngressController
metadata:
```

```

name: default
namespace: openshift-ingress-operator
spec:
  logging:
    access:
      destination:
        type: Container

```

Apply this manifest. This creates a new container named `logs` in the controller's pod. Access log messages are stored there.

```

$ oc project openshift-ingress
$ oc get pod
$ oc describe pod/router-default-6f855cccdb-8vz92 | more

```

Under `Containers`, there is a `router`, which is the main ingress container. View the environment. Interesting details: Certificate information, canonical hostname, load balance algorithm, several volume mounts. What are the mounts for?

The second container is named `logs`. Which command does it run? What are the mounts for?

View the access log

```
$ oc logs -f -c logs router-default-6f855cccdb-8vz92
```

There is considerable traffic, seemingly from internal OpenShift components. Leave this running.

Now generate traffic to the *parksmap* application. Open a separate terminal window and get the URL from the command line:

```
$ oc get route -n devproject
```

NAME	HOST/PORT	PATH	SERVICES	PORT
route-to-parksmap	route-to-parksmap-devproject.apps-crc.testing	/	parksmap-katacoda	8080-tcp

Use curl to access the URL repeatedly:

```

$ while true
do curl -sL route-to-parksmap-devproject.apps-crc.testing >/dev/null
sleep 1
done

```

Watch the ingress log. You should see parksmap entries appear in the log.

Check what happens in the logs container

Launch a shell and look around.

```

$ oc exec -it router-default-6f855cccdb-8vz92 -c logs -- bash
bash-4.4$ ps -ef
UID          PID  PPID  C STIME TTY          TIME CMD
1000610+      1    0   0 04:05 ?        00:00:00 /sbin/rsyslogd -n -i /tmp/rsyslog.pid -f /etc/rsyslog/rsyslog.conf

```

Check the syslog configuration.

```

bash-4.4$ more /etc/rsyslog/rsyslog.conf
$ModLoad imuxsock
$SystemLogSocketName /var/lib/rsyslog/rsyslog.sock
$ModLoad omstdout.so
*. * :omstdout:

```

This means that the log output goes to a socket, not a file. No storage space is used for the logs.

Ingress operator - route admission

The domain name of a route must normally be unique. However, identical domain names in different projects can be permitted by enabling *route admission*.

Confirm that identical domain names are not permitted

First, try to create two routes with the same hostname. This will fail, as expected.

Log on to the web console as *developer*. Create a second project named `project2`. Switch to the Administrator role.

In *devproject*, click **Networking**→**Services** and create a dummy service. You will be presented with a YAML file; just change the name under *metadata* to **dummyservice**.

Then create a route. Name it **dummyroute**, and set the Hostname to **dummy.apps-crc.testing**.

Now switch to `project2` and do the same. You will be able to create the service, but the route will have a status of *rejected*. On the end of the route page, you will find the reason for the rejection: "a route in another namespace holds dummy.apps-crc.testing and is older than dummyroute".

Configure route admission

As kubeadmin, apply the manifest named `ingress-route-admission.yaml`:

```
apiVersion: operator.openshift.io/v1
kind: IngressController
metadata:
  name: default
  namespace: openshift-ingress-operator
spec:
  routeAdmission:
    namespaceOwnership: InterNamespaceAllowed
```

Confirm that identical domain names are possible

Go back to the web console. Remove `dummyroute` and recreate it with the same parameters as before. This will now succeed.

Reset route admission

As kubeadmin, apply the manifest named `ingress-route-admission-off.yaml`:

```
apiVersion: operator.openshift.io/v1
kind: IngressController
metadata:
  name: default
  namespace: openshift-ingress-operator
spec:
  routeAdmission:
    namespaceOwnership: Strict
```

Network Policies

Prepare the pods

You will create Network Policies between pods in one project and between two projects.

Add two applications in the web console, both using the NGINX template. You will find it under **Developer Catalog** → **All Services** → **Templates** → **Other** on the **+Add** page. Name one group **"green"**, the other **"blue"**. The pods will have a label **"name: blue"** or **"name: green"**.

Scale both up to 2 (not higher - avoid memory problems).

On the command line, test success with

```
$ oc get pod -l name=blue
$ oc get pod -l name=green
```

Create a Network Policy for traffic between the pod groups

Manifest `np-onlyblue.yaml` is effective for blue pods (see `podSelector`) and only allows traffic from blue pods (ingress from `podSelector`).

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: only-blue
  namespace: devproject
spec:
  ingress:
    - from:
        - podSelector:
            matchLabels:
              name: blue
  podSelector:
    matchLabels:
      name: blue
  policyTypes:
    - Ingress
```

Apply the manifest.

Test the policy

List all blue and green pods with option `-o wide` to print their IP addresses.

```
$ oc get pod -o wide -l name=blue
$ oc get pod -o wide -l name=green
```

Enter a blue pod and run `curl` with an IP address from another blue pod. You should get an error message *Connection Refused* which indicates that the connection request was received by the other side. Also use the IP address of a green pod. The result should be identical.

Enter a green pod and again run `curl` with an IP address from a **blue** pod. This command should hang; no error message. This shows that the blue pods are protected from outside access.

Create a Network Policy for traffic between projects

Create a project named **project2** (if it doesn't exist). In this project, launch another NGINX app with one pod, called **red**.

Enter the red pod and try to `curl` to a blue and a green pod. The green pods should be accessible (although they refuse connection), the blue pods not.

`np-allowp2.yaml` allows access to the blue pods from project2. Apply it and test again if you can curl to a blue pod from a red pod.

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: from-project2
  namespace: devproject
spec:
  ingress:
    - from:
        - namespaceSelector:
            matchLabels:
              project: project2
        - podSelector: {}
  podSelector:
    matchLabels:
      name: blue
  policyTypes:
    - Ingress
```

The highlighted podSelector line selects an empty set of pods. Somewhat unintuitively, this means all pods.